Addendum to SemVer: System-on-Chip Designs

Overview

SoC (System-on-Chip) designs are similar to sofware in many ways, e.g. they are written in human-readable computer languages and often developed using version control, but differ in what comprises an "API". This is an addendum to the SemVer 2.0.0 specification which clarifies how to apply SemVer to SoC designs. It is assumed that you have read and understood:

- RFC 2119
- SemVer 2.0.0

The main difference between software and SoC designs is in what constitutes a public Application Programming Interface (API).

In software libraries written in the C language, an API includes:

- Names of header files which a library user can reference via `#include`.
- Names of preprocessor macros (defined with the `#define` directive).
- Semantics of preprocessor macros.
- Values of constants.
- Names of exposed functions.
- Order of arguments to exposed functions.
- Anything else which a library user might reasonably rely on.

In command-line applications, an API includes:

- Names of command-line options, flags, and sub-commands.
- Default values of command-line options.
- Precedence of configuration files.
- Anything else which an application user might reasonably rely on.

Those examples of software public APIs demonstrate that an API can be described more generally as *anything which a user might reasonably rely on.*

SoC designs are typically written in specialised languages for digital logic such as Verilog (IEEE1364), SystemVerilog (IEEE1800), or VHDL (IEEE1076) which facilitate synthesis to physical digital logic circuits. Different from software, SoC designs have users with fundamentally different requirements related to higher-level designs, high-level software, and (most critically) physical implementation. These downstream users likely have differing perspectives about what constitutes the most important part of the public API – A system-level software user might depend on the address and reset value of a register, but not depend on the hierarchical path to the corresponding FF (flip-flop) because that doesn't affect their software. In contrast, a user working on physical implementation might view that register address as a minor detail, but depend on the hierarchical path of the FF to ensure that it is implemented with the correct type of cell.

Downstream Users and Auxiliary Components

Downstream projects, i.e. those which depend on your SoC design(s), often fall into these categories:

1. Documentation: Describe the intention, features, and operation of your design as required by other users. The extent of overlap with other categories is highly specific to each particular project.
2. Integration: Include your design as a hierarchical component in a larger system. In SystemVerilog, this means declaring an instance of your `module` with suitable connections to parameter and signal ports.
3. Verification: Check that your design meets specifications. Includes both simulation and formal methods, and checks may be written in a different language from your SoC design, e.g. TCL, SystemC, or Python.
4. Implementation: Convert your abstract design into a concrete realisation on a physical ASIC or FPGA platform. Includes synthesis, layout, and modifications for testability.
5. Software: Most modern SoC designs feature some software-programmable component, which will view your SoC design from a memory-mapped perspective. Includes system-level tests (overlapping with verification), validation and characterisation (overlapping with implementation), firmware, and possibly end-user applications.

In addition to the main description of digital logic, e.g. a collection of SystemVerilog files, a SoC design will usually include auxiliary components which may (or may not) correspond to additional files.

- Filesystem structure of the release delivery and associated filelists.
- Standards which the source code adheres to. Includes standards of all types from file encoding (e.g. ASCII vs UTF8) and whitespace/formatting rules to the specification language (e.g. Verilog vs SystemVerilog) and naming conventions. Other users may have flows which extract information which depend on these seemingly minor details.
- Hierarchical paths and specific identifiers. Used by verification and implementation to identify and manipulate parts of the design. In verification, specific FFs may be forced to obtain coverage on a difficult-to-reach state. In implementation, specific FFs may be selected as requiring special treatment to meet timing.
- Constraints on clocks, timing, pin-mapping, cell placement, routing, etc.
- Waivers on errors or warnings from particular tools.
- Power intent, normally specified with UPF (IEEE1801).
- Abstract models and unittests. Other users (likely excluding implementation) may depend on these to validate their own work.
- Address map and structure of software-visible registers. All types of software may depend on specific addresses, the layout of register fields, and more subtle attributes like reset values and volatility.

When releasing a new SoC design version, it is important to consider how your changes all components of the release affect all users developing downstream projects.

Changes in SystemVerilog

An illustrative example, shown in SystemVerilog, is useful to demonstrate API
components of a typical SoC peripheral where sequential logic is implemented
with D-type FFs. Let's say that our module `Alu` performs arithmetic operations
on its inputs, drives known values on its outputs, and provides register access
via the APB protocol. In the most recently released version, there is one
configuration register called `CFG` at the address `12'h444`, with a reset value of
`32'd5`, arranged as two fields `CFG[2:1]=OPERATION` and `CFG[0]=ENABLE`.

```systemverilog
module Alu
  #(parameter int RESULT_W = 16
  )
  ( input  var logic [1:0][7:0]    i_operands
  , output var logic [RESULT_W-1:0] o_resultant
  , APB.slave                      ifc_APB
  );

  localparam bit MYCONSTANT = 1'b1;

  // Combinatorial assignment via `always_comb`, `assign`,
  // or connection to sub-module.
  logic foo_d;

  // Sequential assignment via `always_ff`.
  logic foo_q;

  // ... snip ...

  ArithmeticPipe u_pipeA1
    ( .i_opA  (i_operands[3:0])
    , .i_opB  (i_operands[7:4])
    , .o_taps (foo_d)
    );

  always_ff @(posedge ifc_APB.clk, posedge ifc_APB.arst)
    if (ifc_APB.arst)
      foo_q <= '0;
    else if (updateFoo)
      foo_q <= foo_d;

  // ... snip ...

endmodule
```

The public API of this module may consist of the module declaration, APB reg-
isters, hierarchical paths to sequential elements, and other packaged components
like helper scripts and design constraints.

MAJOR Versions

Given a version number MAJOR.MINOR.PATCH, increment the:

1. MAJOR version when you make incompatible API changes

Referencing the example, the MAJOR version must be incremented with any of the following changes:

1. Modified module name which integrators use to declare an instance of the peripheral, e.g. `Alu` → `MyArithmetic`. Existing code using the name `Alu` will not elaborate unchanged.
2. Removed parameter port, e.g. ~~`RESULT_W`~~. Existing code overriding the parameter value will not elaborate unchanged.
3. Modified parameter port kind, e.g. `parameter` → `localparam`, i.e. overridable to non-overridable. Existing code overriding the parameter value will not elaborate unchanged.
4. Modified parameter port name, e.g. `RESULT_W` → `OUT_WIDTH`. Existing code using the name `RESULT_W` will not elaborate unchanged.
5. Modified parameter port default value, e.g. `16` → `5`, including addition or removal of the explicit default value. Existing code may depend on the default value for critical functionality.
6. Removed signal port, e.g. ~~`o_resultant`~~. Existing code using that port will not elaborate unchanged.
7. Modified signal port datatype, e.g. `logic [1:0][7:0]` → `logic [15:0]`. Existing code may depend on the size and structure of the port datatype, and input expressions may be cast to an unexpected width or datatype.
8. Modified signal port name, e.g. `i_operands` → `i_numbers`. Existing code using the name `i_operands` will not elaborate unchanged.
9. Removed interface port, e.g. ~~`ifc_APB`~~. Existing code using the APB interface will not elaborate unchanged.
10. Modified interface port type, e.g. `APB.slave` → `AXI.slave`. Existing code using the APB interface will not elaborate unchanged.
11. Modified interface port name, e.g. `ifc_APB` → `myApb`. Existing code using the name `ifc_APB` will not elaborate unchanged.
12. Removed or modified sequential signal name, e.g. `foo_q` → `bar_q`. Existing code referencing `foo_q` will not find the inferred FF(s). You may not notice the breakage until your colleagues in physical implementation notify you that their scripts don't work. In the worst cases, FFs requiring special treatment can be silently missed.
13. Any added, removed, or renamed hierarchical middle layer, e.g. `Alu.u_pipe` → `Alu.u_wrapperA.u_pipe`. Existing code, particularly for physical implementation, may depend on the hierarchical names including generate blocks.
14. Removed, or renamed hierarchical bottom layer, e.g. `Alu.u_pipe1` → `Alu.u_pipe[1]`. Existing code, particularly for physical implementation, may depend on the hierarchical names including generate loops.
15. Added, removed, or modified any machine-readable comment, e.g. tool-specific directives like `// synopsys parallel_case`. Existing flows are likely to depend on these for critical functionality.
16. Removed software-accessible register, e.g. ~~`CFG`~~. Existing system software

accessing the `CFG` address will not operate equivalently.

17. Modified software-accessible register address, e.g. `12'h444` → `12'h888`. Existing system software accessing the address `0x444` will not operate equivalently.
18. Modified software-accessible register field layout, e.g. `CFG[0]=ENABLE` → `CFG[31]=ENABLE`. Existing system software accessing the register will not operate equivalently.
19. Modified software-accessible register reset value, e.g. `32'd5` → `32'd0`. Existing system software accessing the register will not operate equivalently, particularly software performing non-atomic read-modify-write operations on startup like `cfg->operation++`.

To summarise, the MAJOR version must be incremented with any changes which *require* updates to any projects that fetch the newly released version. Note, changes to match the documentation of a previous release should be considered bug fixes, so may only warrant a MINOR increment.

MINOR Versions

Given a version number MAJOR.MINOR.PATCH, increment the:

2. MINOR version when you add functionality in a backwards compatible manner

Where SemVer specifies *adding* functionality, SoC designs must update at least the MINOR version with any of the following modifications:

1. Added parameter port, e.g. `ANOTHER`. Existing code will elaborate unchanged.
2. Modified parameter port datatype, e.g. `int` to `bit [3:0]`, including removal of the explicit datatype. Existing code may elaborate unchanged, but override values may be cast to an unexpected width or datatype. If existing code needs changes to elaborate with the updated version, then increment MAJOR instead.
3. Added signal port, e.g. `output o_another`. Existing code may elaborate unchanged and a new signal port implies new functionality.
4. Modified signal port direction, e.g. `myport` → `output myport`. Default direction is `inout`. Existing code may elaborate unchanged, but simulation semantics may be different. If existing code needs changes to elaborate with the updated version, then increment MAJOR instead.
5. Modified signal port nettype, e.g. `input logic` → `input var logic`. Default nettype of `input` and `inout` signal ports with datatype `logic` is `tri`, but for `output` ports it's `var`. Existing code may elaborate unchanged, but simulation semantics may be different. If existing code needs changes to elaborate with the updated version, then increment MAJOR instead.
6. Added interface port, e.g. `OCP.slave`. Existing code may elaborate unchanged and a new interface port implies new functionality.
7. Modified sequential signal datatype or expression, e.g. `logic [1:0] foo_q` → `FooEnum_t foo_q`. Backwards-compatible changes only require a MINOR increment, but incompatible changes like reducing the *intended* width of a FF vector require a MAJOR increment.

8. Added hierarchical bottom layer, e.g. `Alu.u_pipeA2`. New hierarchy implies new functionality, not just a bug fix.
9. Added software-accessible register, e.g. `STATUS`. Existing system software will not operate equivalently, and updated software may use the new functionality.

To summarise, the MINOR version must be incremented with any changes which add or modify functionality in a manner which *does not require* downstream users to make changes. If downstream users are required to make changes to their project in order to accept the new version, increment MAJOR instead.

PATCH Versions

Given a version number MAJOR.MINOR.PATCH, increment the:

3. PATCH version when you make backwards compatible bug fixes

As with SemVer, only backwards-compatible changes (for all downstream users) are allowed within a PATCH increment version.

1. Added, removed, or modified internal constant, e.g. `MYCONSTANT` → `BETTERNAME`. Internal constants should not be relied upon downstream.
2. Added, removed, or modified internal combinational signal, e.g. `foo_d` → `bar_d`. Internal combinational signals should not be relied upon downstream. Exemption: If you change signals which are *intended* to be probed or forced by downstream users, increment MAJOR instead, e.g. `disableChecks` → `turnOffChecks`.
3. Added internal sequential signal, e.g. `new_q`. Additional FFs will affect area, power, achieveable fmax and cost, but are unlikely to break physical implementation flows outright. Note, removed or renamed internal signals require a MAJOR increment.
4. Any machine-readable status tracker comment, e.g. `/* TODO: Something */`. Note, if there are updated status tracker comments, there's a good chance the changes also involve enough to warrant a MINOR or MAJOR increment.
5. Any human-only comment, e.g. `/* Isn't this nice */`.

Exemptions

The public API is restricted to the reasonable ways that users are expected to use your release. In-house projects may use this restriction to avoid incrementing MAJOR too often, i.e. the distinction between a breaking change and a bugfix can be redefined if you (1) *identify **all** downstream projects/users* and (2) *obtain explicit agreement from **all** users*. This exemption allows larger SoC designs like subsystem and chip-level projects to make arbitrary changes under MINOR increments while reserving MAJOR increments for project-specific milestones. Only in-house projects may use this exemption because publicly available projects cannot identify all downstream projects/users.

SemVer for SystemVerilog Cheatsheet

How	What	Increment
`mod`	Top-level module name	MAJOR
`add`	Parameter port	MINOR
`rem`	Parameter port	MAJOR
`mod`	Parameter port kind	MAJOR
`mod`	Parameter port datatype	MINOR
`mod`	Parameter port name	MAJOR
`mod`	Parameter port default value	MAJOR
`add`	Signal port	MINOR
`rem`	Signal port	MAJOR
`mod`	Signal port direction	MINOR
`mod`	Signal port nettype	MINOR
`mod`	Signal port datatype	MAJOR
`mod`	Signal port name	MAJOR
`add`	Interface port	MINOR
`rem`	Interface port	MAJOR
`mod`	Interface port type	MAJOR
`mod`	Interface port name	MAJOR
any	Internal constant	PATCH
any	Combinatorial signal	PATCH
`add`	Sequential signal	PATCH
`rem`	Sequential signal	MAJOR
`mod`	Sequential signal name	MAJOR
`mod`	Sequential signal datatype	MINOR
`mod`	Sequential signal expression	MINOR
any	Hierarchy middle layer	MAJOR
`add`	Hierarchy bottom layer	MINOR
`mod`	Hierarchy bottom layer	MAJOR
`rem`	Hierarchy bottom layer	MAJOR
any	Tool directive comment	MAJOR
any	Status tracker comment	PATCH
any	Human-only comment	PATCH
`add`	Software register	MINOR
`rem`	Software register	MAJOR
`mod`	Software register address	MAJOR
`mod`	Software register field layout	MAJOR
`mod`	Software register reset value	MAJOR